

Task 1: Corner Detection

1.2, 1.3

```
1 def conv2(img, conv_filter):
2     # flip the filter
3     f_siz_1, f_size_2 = conv_filter.shape
4     conv_filter = conv_filter[range(f_siz_1 - 1, -1, -1),
5                               :, range(f_siz_1 - 1, -1, -1)]
6     pad = (conv_filter.shape[0] - 1) // 2
7     result = np.zeros((img.shape))
8     img = np.pad(img, pad_width=pad, mode='constant',
9                  constant_values=0)
10    filter_size = conv_filter.shape[0]
11    for r in np.arange(img.shape[0] - filter_size + 1):
12        for c in np.arange(img.shape[1] - filter_size +
13                           1):
14            curr_region = img[r:r + filter_size, c:c +
15                               filter_size]
16            curr_result = curr_region * conv_filter
17            conv_sum = np.sum(curr_result) # Summing the
18            # result of multiplication.
19            result[r, c] = conv_sum # Saving the
20            # summation in the convolution layer feature map.
```

```

21     Create gaussian filter based on specified parameter,
22     shape and sigma.
23
24     The 2DGaussian will be a function of distance from
25     the center and the provided sigma.
26
27
28     P ~ Norm(0, sigma)
29     h[i, j] = P(d) where P(d) is the probability density
30     function of the specified gaussian distribution,
31     and d is distance from center (0,0) to the position
32     i, j.
33
34     The closer d is to 0, the higher h[i, j] is.
35     The output matrix (kernel) will be normalized to have
36     sum of 1
37     """
38
39
40
41
42 # Parameters, add more if needed
43 sigma = 5
44 # thresh = 0.01
45 harris_kernel_shape = (5, 5)
46 k = 0.05
47
48 # Derivative masks
49 dx = np.array([[-1, 0, 1], [-1, 0, 1], [-1, 0, 1]])
50 dy = dx.transpose()
51
52 try:
53     bw = plt.imread('0.png')
54 except:
55     bw = plt.imread('Task1/Harris-1.jpg')
56     bw = cv2.cvtColor(bw, cv2.COLOR_RGB2GRAY)

```

```

57
58
59 bw = np.array(bw * 255, dtype=int)
60
61 # computer x and y derivatives of image
62 Ix = conv2(bw, dx)
63 Iy = conv2(bw, dy)
64
65 # gaussian filter
66 g = fspecial(
67     (
68         max(1, np.floor(3 * sigma) * 2 + 1),
69         max(1, np.floor(3 * sigma) * 2 + 1)
70     ),
71     sigma
72 )
73
74 Iy2 = conv2(np.power(Iy, 2), g)
75 Ix2 = conv2(np.power(Ix, 2), g)
76 Ixy = conv2(Ix * Iy, g)
77
78 ######
79 # Task: Compute the Harris Cornerness
80 #####
81
82 # if w(x, y) is (1, 1), don't need to apply convolution
83 if harris_kernel_shape == (1, 1):
84     lambda_1 = Ix2
85     lambda_2 = Iy2
86 # else, for each pixel position, we will add up Ix2 and
87 # Iy2 of each pixel that overlaps with the filter
88 else:
89     kernel = np.ones(harris_kernel_shape)
90     lambda_1 = conv2(Ix2, kernel)
91     lambda_2 = conv2(Iy2, kernel)
92
93 # calculate determinant and trace of matrix M for each
94 # pixel in the image, and finally compute R

```

```

93 det_m = lambda_1 * lambda_2 - np.square(conv2(Ixy,
94 kernel))
95 trace_m = lambda_1 + lambda_2
96 R = det_m - (k * np.square(trace_m))
97 #####
98 ##### # Task: Perform non-maximum suppression and
99 ##### # thresholding, return the N corner points
100 ##### # as an Nx2 matrix of x and y coordinates
101 #####
102 #####
103 def non_maximum_suppression(R):
104     """
105         Perform non-maximum suppression by finding pixels
106         that have intensity greater than 8 neighbors.
107
108         input:
109             R: 2-d numpy array of corner response values
110         returns:
111             corners: list of tuples indicating corners
112             detected
113             arr: 2-d numpy binary array where arr[i, j] = 1
114             indicates corner as detected by non-maximum suppression
115             """
116
117         corners = []
118         for i in range(1, len(R)-1):
119             for j in range(1, len(R[0])-1):
120                 # check if R[i,j] is a corner by comparing
121                 # each pixel to 8 neighbors
122                 # collect corners
123                 if (
124                     (R[i,j] >= R[i+1,j]) & (R[i,j] >= R[i-
125                     1,j]) & (R[i,j] >= R[i,j+1]) & (R[i,j] >= R[i,j-1]) &
126                     (R[i,j] >= R[i+1,j+1]) & (R[i,j] >= R[i-
127                     1,j-1]) & (R[i,j] >= R[i-1,j+1]) & (R[i,j] >= R[i-1,j-1])
128                 ):
129                     corners += (i, j),
130
131

```

```

125     # create a binary array from corners array to
126     # visualize corners
127     arr = np.zeros(R.shape)
128     for corner in corners:
129         arr[corner[0], corner[1]] = 1
130     return corners, arr
131 corners, Nx2 = non_maximum_suppression(R)

```

1.4

Fixed parameters

Fixed parameters: sigma=5, harris_kernel_width=5, k=0.05

Harris-1.jpg 239 corners



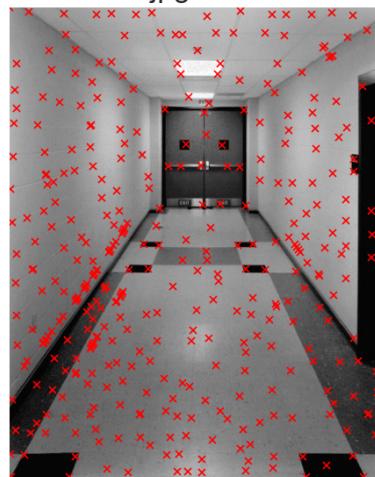
Harris-2.jpg 314 corners



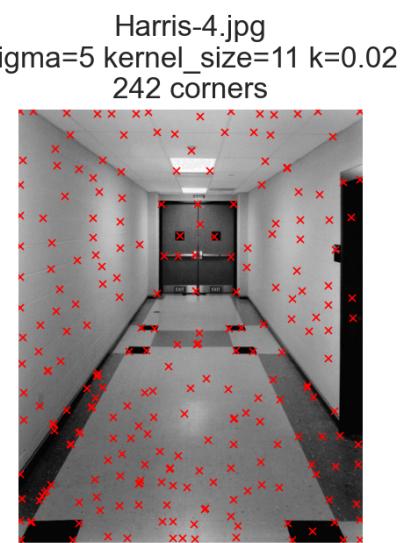
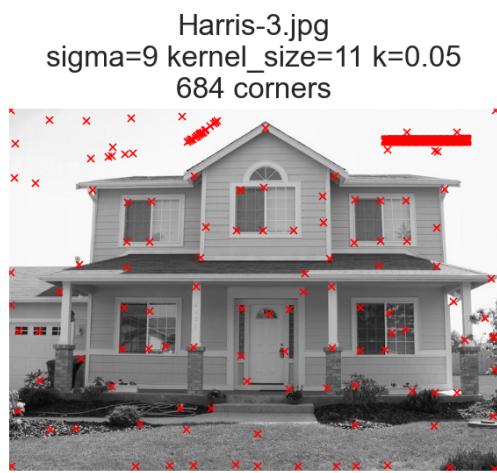
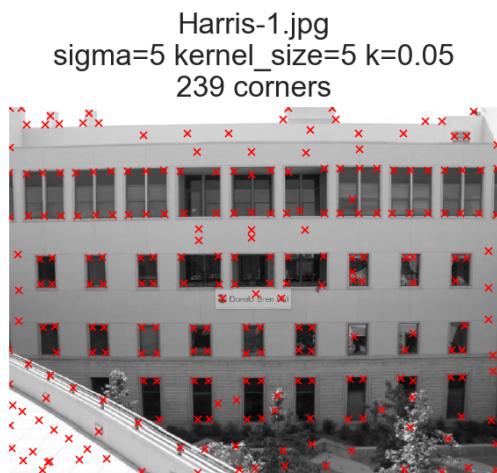
Harris-3.jpg 6290 corners



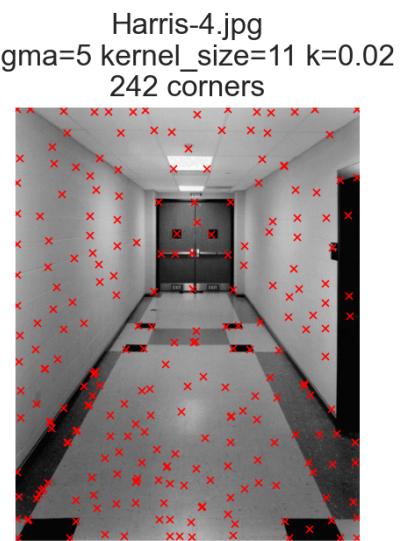
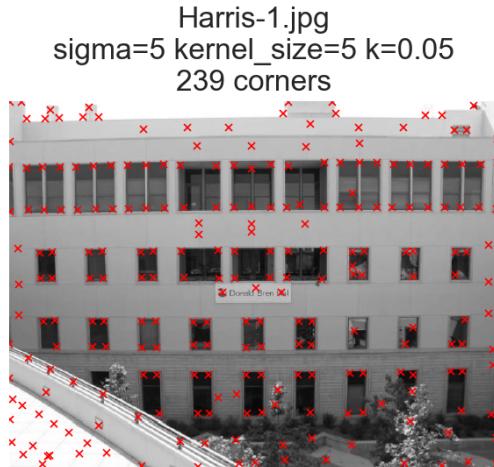
Harris-4.jpg 374 corners



Variable parameters



1.5. Compare results with `cv2.cornerHarris()`



Factors that affect the performance of Harris corner detection

Harris corner detection is invariant to location, but not scale. As such, the performance will depend on the degree of gaussian smoothing determined by `sigma`, the kernel size of the sobel filter, the size of neighborhood considered when determining whether `img[i][j]` is a corner or not, and the `k` value in calculating the cornerness score R .

1.6

Harris-5.jpg

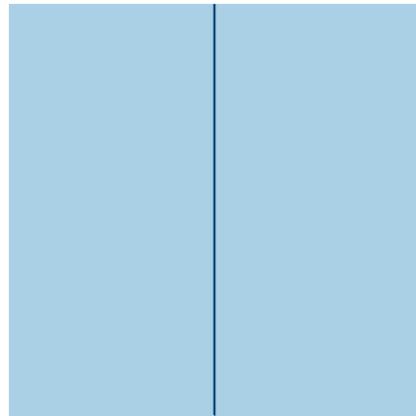


253001 corners



Analyse the results why we cannot get corners by discussing and visualising your corner response scores of the image.

Cornerness R visualized using `plt.imshow()` with `cmap RdBu_r` (blue is lower).



R value is approximately 0 for almost all pixel positions in the image. The only pixels with high R scores are at top and bottom left corners, and top and bottom center because the right half of the image is all 0s. Corners were detected on the left corners only because of the 0 padding of width 1 that's applied to the image in `conv2()` function. Corners were also detected in the top and bottom center because of the boundary between the white and black regions, in addition to the 0 padding. Aside from the influence from 0 padding, no corners were detected because there are no corners in the image.

```
1 Top left corner
2 [[1.95029123 2.09187119 0.11344908]
3 [2.09187119 2.41787998 0.51097243]
4 [0.11344908 0.51097243 0.17606731]]
5
6 Top right corner
7 [[0. 0. 0.]]
```

```

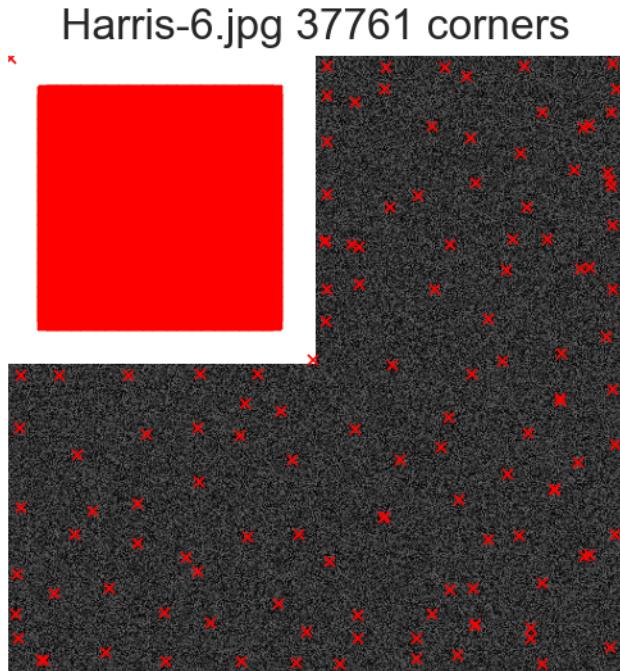
8 [ 0. 0. 0. ]
9 [ 0. 0. 0. ]]
10
11 Bottom left corner
12 [[ 0.11344908 0.51097243 0.17606731]
13 [ 2.09187119 2.41787998 0.51097243]
14 [ 1.95029123 2.09187119 0.11344908]]
15
16 Bottom right corner
17 [[ 0. 0. 0. ]
18 [ 0. 0. 0. ]
19 [ 0. 0. 0. ]]
20
21 Top center
22 [[ 0.17165154 2.65415093 3.26432971 0.96915069
0.06661764]
23 [ 0.57056862 2.97199509 3.20967762 0.64287179
-0.0273949 ]
24 [ 0.19054354 0.57293854 -0.33744827 -1.04884253
-0.27012009]]
25
26 Bottom center
27 [[ 0.19054354 0.57293854 -0.33744827 -1.04884253
-0.27012009]
28 [ 0.57056862 2.97199509 3.20967762 0.64287179
-0.0273949 ]
29 [ 0.17165154 2.65415093 3.26432971 0.96915069
0.06661764]]

```

Harris-5Top & Bottom center R scores



1.7



Task 2: KMeans

2.1 KMeans

```
1  class KMeans:
2      def __init__(self, X: np.ndarray, k: int):
3          """initialize parameters and randomly create
4          centroids in the domain of input data"""
5          self.X = X
6          self.k = k
7          self.centroids = np.random.uniform(size=(self.k,
8                                              self.X.shape[1]))
9          self.pred = np.zeros(X.shape[0])
10
11     @staticmethod
12     def euclidean_distance(X, centroids):
13         """calculate pair-wise euclidean distance"""
14         k = centroids.shape[0]
15
16         diff = X[np.newaxis,:,:] -
```

```

15         assert diff.shape == (k, X.shape[0], X.shape[1])
16
17         dist = np.sqrt(np.sum(diff**2, axis=-1))
18         dist = dist.transpose()
19         assert dist.shape == (X.shape[0], k)
20
21     return dist
22
23 def predict(self, verbose=False):
24     i = 1
25     start_time = time.time()
26
27     while True:
28         # calculate distance between every point and
29         # every centroid
30         dist = KMeans.euclidean_distance(self.X,
31                                         self.centroids)
32
33         # assign each point to a class with nearest
34         # centroid
35         pred = np.argmin(dist, axis=1)
36
37         # calculate new centroid and wss for each
38         # cluster
39         new_centroids = []
40         wss = []
41         for c in range(self.k):
42             if len(self.X[pred==c]) == 0:
43                 new_centroid =
44                     np.random.uniform(size=self.centroids.shape[1])
45             else:
46                 new_centroid =
47                     np.mean(self.X[pred==c], axis=0)
48                 new_centroids += new_centroid,
49
50             # within-cluster sum of squared difference
51             wss += np.sum((self.X[pred==c] -
52                           new_centroid)** 2),
53
54         new_centroids = np.array(new_centroids)

```

```

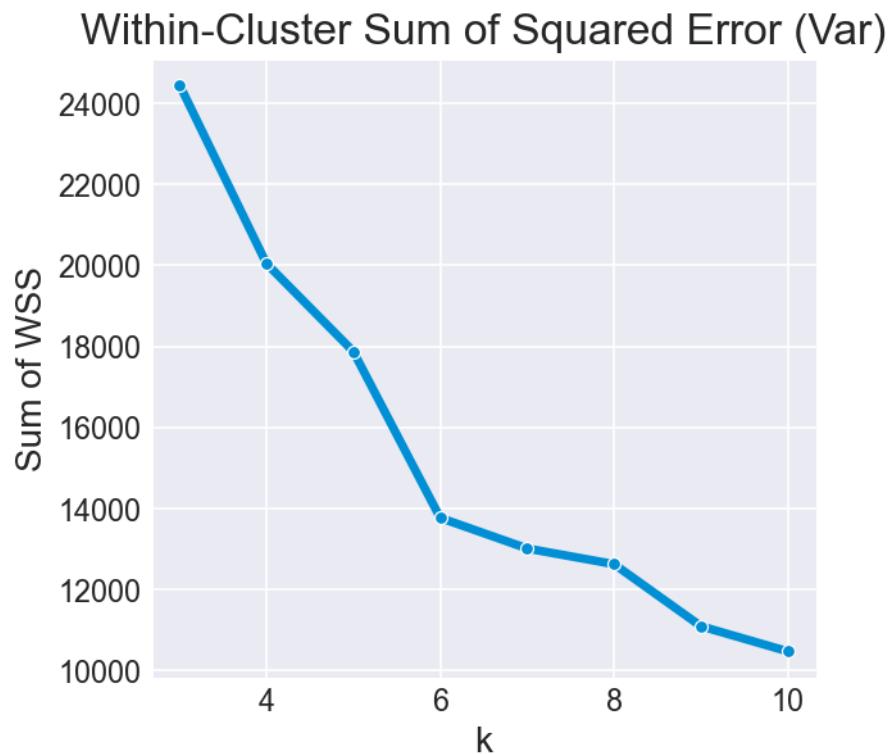
48         assert self.centroids.shape ==
new_centroids.shape
49         self.centroids = new_centroids
50         self.wss = np.array(wss)
51
52     if all(pred == self.pred):
53         return time.time() - start_time
54     else:
55         if i % 5 == 0 and verbose:
56             elapsed = time.time() - start_time
57             print(f'iteration: {i:>3}\twss:
{self.wss.sum():^10.2f}\ttime in seconds:
{elapsed:>5.2f}')
58
59         last_wss = self.wss
60         self.pred = pred
61         i += 1
62
63 def my_kmeans(data, start, end):
64 """
65     Run KMeans for various k values in the given range and
return the scores and time to converge
66 """
67     scores = []
68     models = []
69     conv_time = []
70
71     assert start < end
72
73     for k in range(start, end):
74         kmeans = KMeans(data, k=k)
75         t_elapsed = kmeans.predict()
76
77         wss = kmeans.wss.sum()
78         scores += wss,
79         models += kmeans,
80         conv_time += t_elapsed,
81     return np.array(scores), conv_time

```

2.2

With pixel coordinates

Image 1: M&M



Using $k=8$

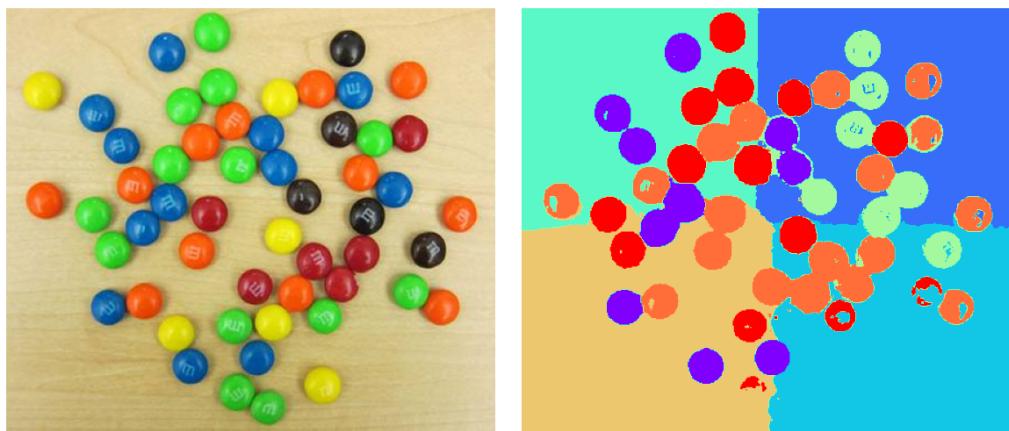
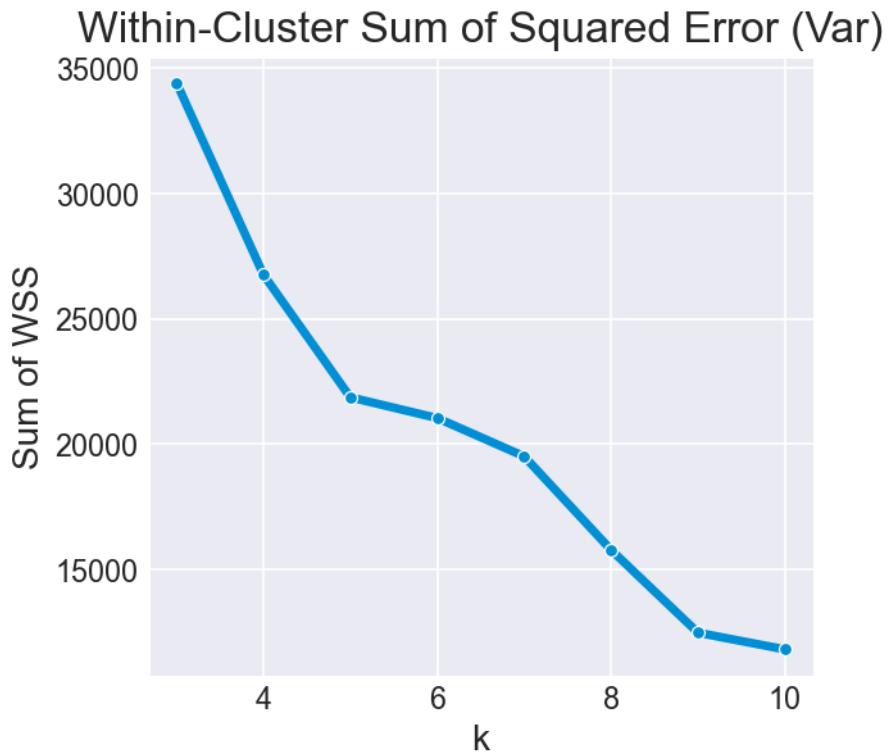
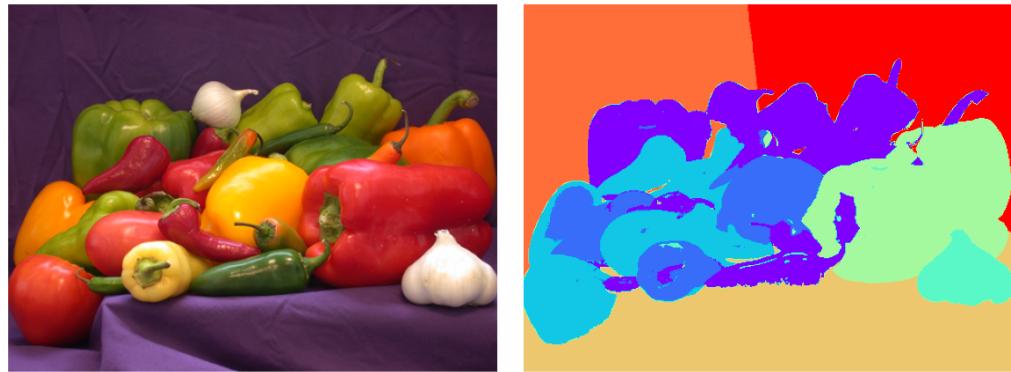


Image 2: Peppers



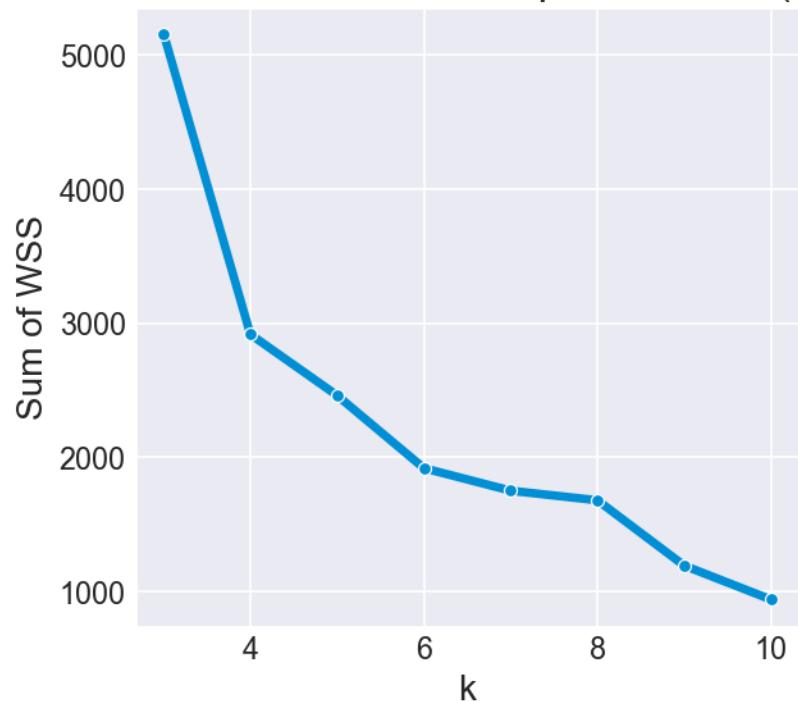
Using $k=8$



Without pixel coordinates

Image 1: M&M

Within-Cluster Sum of Squared Error (Var)

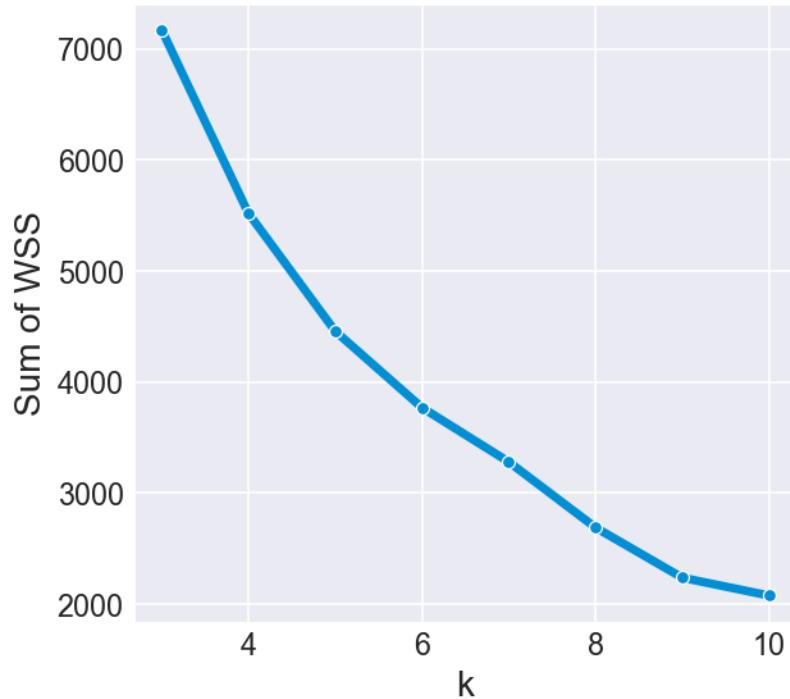


Using $k=7$

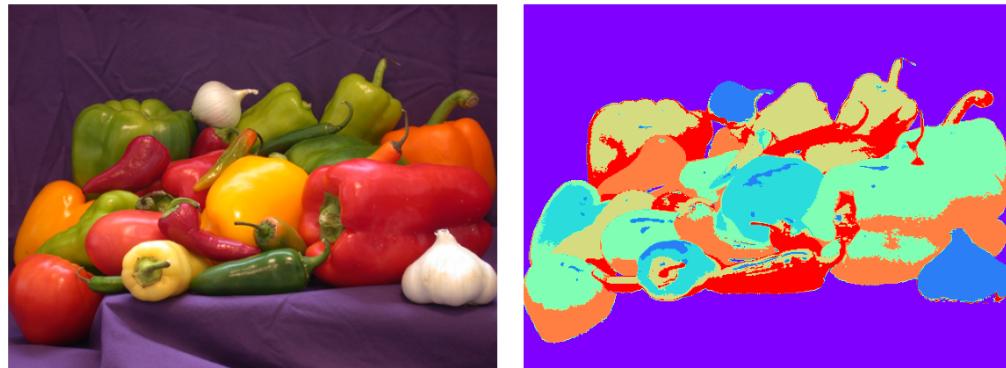


Image 2: Peppers

Within-Cluster Sum of Squared Error (Var)



Using $k=7$



2.3 KMeans++

Key differences

The key difference between (Naive) KMeans and KMeans++ is in the initialization of the centroids.

While KMeans randomly initializes centroids in the input data space using uniform distribution, KMeans++ randomly picks one data point as the first centroid, and picks the rest of $k-1$ centroids by choosing the data point that is most distant from the closest centroid.

```
1 | class KMeansPP:
```

```

2     def __init__(self, X: np.ndarray, k: int):
3         self.X = X
4         self.k = k
5         self.initialize_centroids()
6         self.pred = np.zeros(self.X.shape[0])
7
8     @staticmethod
9     def euclidean_distance(X, centroids):
10        """calculate pair-wise euclidean distance"""
11        k = centroids.shape[0]
12
13        diff = X[np.newaxis,:,:] -
14        centroids[:,np.newaxis,:]
15        assert diff.shape == (k, X.shape[0], X.shape[1])
16
17        dist = np.sqrt(np.sum(diff**2, axis=-1))
18        dist = dist.transpose()
19        assert dist.shape == (X.shape[0], k)
20
21        return dist
22
23    def initialize_centroids(self):
24        X = copy.deepcopy(self.X)
25        idx = np.random.randint(low=0, high=X.shape[0])
26        centroid = X[idx]
27        centroids = np.array([centroid])
28        X = np.delete(X, idx, 0)
29
30        for i in range(1, self.k):
31            # calculate distance to the nearest centroid
32            dist = KMeansPP.euclidean_distance(X,
33            centroids)
34            dist_to_nearest_centroid = np.min(dist, 1)
35
36            # use the point furthest from the nearest
37            # cluster as the next centroid
38            idx = np.argmax(dist_to_nearest_centroid ** 2)
39            centroid = X[idx]
40            centroids = np.vstack((centroids, centroid))
41            X = np.delete(X, idx, 0)

```

```

39
40         self.centroids = centroids
41
42     def predict(self, verbose=False):
43         i = 1
44         start_time = time.time()
45
46         while True:
47             # calculate distance between every point and
48             # every centroid
49             dist = KMeans.euclidean_distance(self.X,
50                                              self.centroids)
51
52             # assign each point to a class with nearest
53             # centroid
54             pred = np.argmin(dist, axis=1)
55
56             # calculate new centroid and wss for each
57             # cluster
58             new_centroids = []
59             wss = []
60             for c in range(self.k):
61                 if len(self.X[pred==c]) == 0:
62                     new_centroid =
63                     np.random.uniform(size=self.centroids.shape[1])
64                 else:
65                     new_centroid =
66                     new_mean = np.mean(self.X[pred==c], axis=0)
67                     new_centroids += new_mean,
68
69                     # within-cluster sum of squared difference
70                     wss += np.sum((self.X[pred==c] -
71                                   new_mean) ** 2),
72
73             new_centroids = np.array(new_centroids)
74             assert self.centroids.shape ==
75             new_centroids.shape
76             self.centroids = new_centroids
77             self.wss = np.array(wss)
78
79             if all(pred == self.pred):
80
81

```

```
72         return time.time() - start_time
73     else:
74         if i % 5 == 0 and verbose:
75             elapsed = time.time() - start_time
76             print(f'iteration: {i:>3}\twss:
77 {self.wss.sum():^10.2f}\ttime in seconds:
78 {elapsed:>5.2f}')
79
80
81 def my_kmeanspp(data, start, end):
82     scores = []
83     models = []
84     conv_time = []
85
86     assert start < end
87
88     for k in range(start, end):
89         kmeanspp = KMeansPP(data, k=k)
90         t_elapsed = kmeanspp.predict()
91
92         wss = kmeanspp.wss.sum()
93         scores += wss,
94         models += kmeanspp,
95         conv_time += t_elapsed,
96     return np.array(scores), conv_time
```

KMeans++ Initialized Centroids in CIE-Lab Space

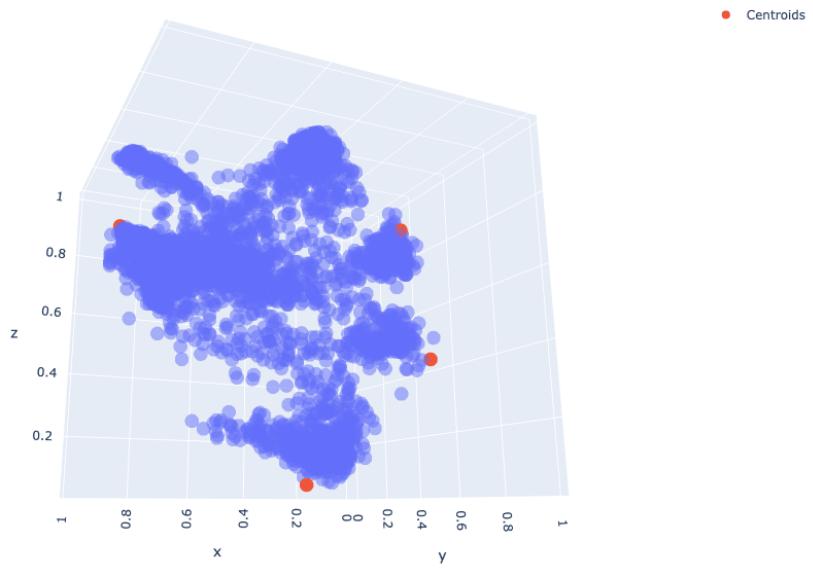
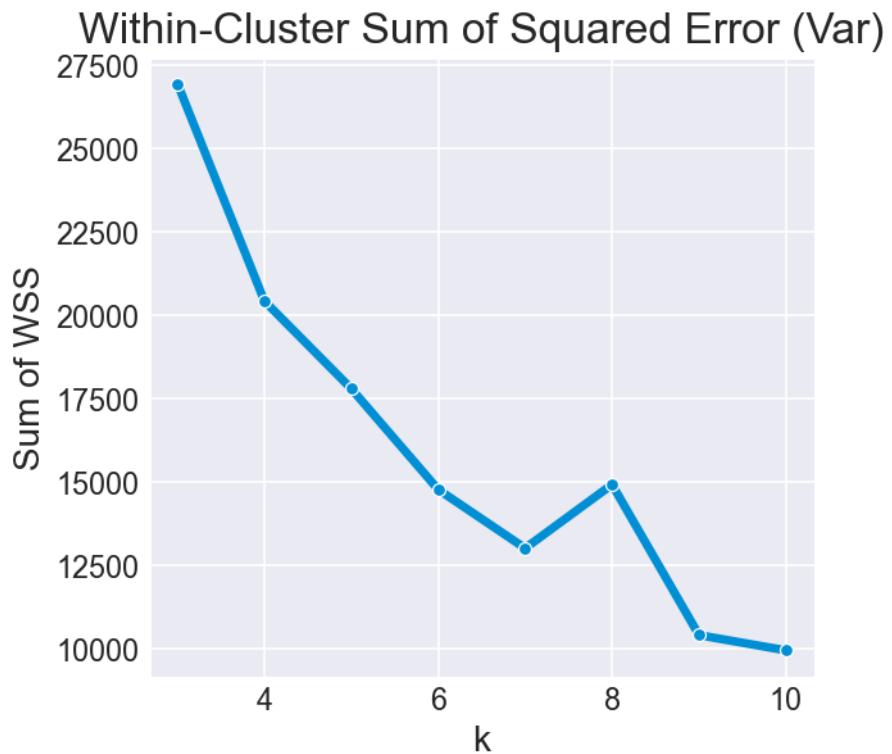


Image 1: M&M



Using $k=7$

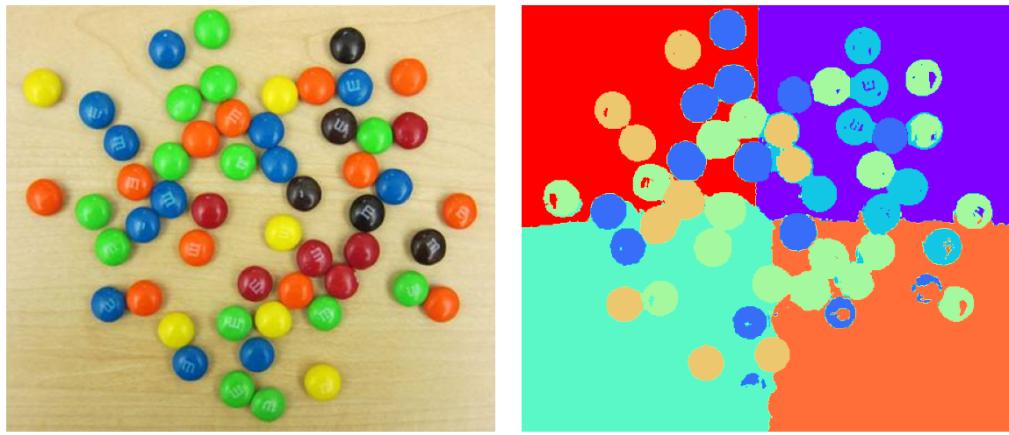
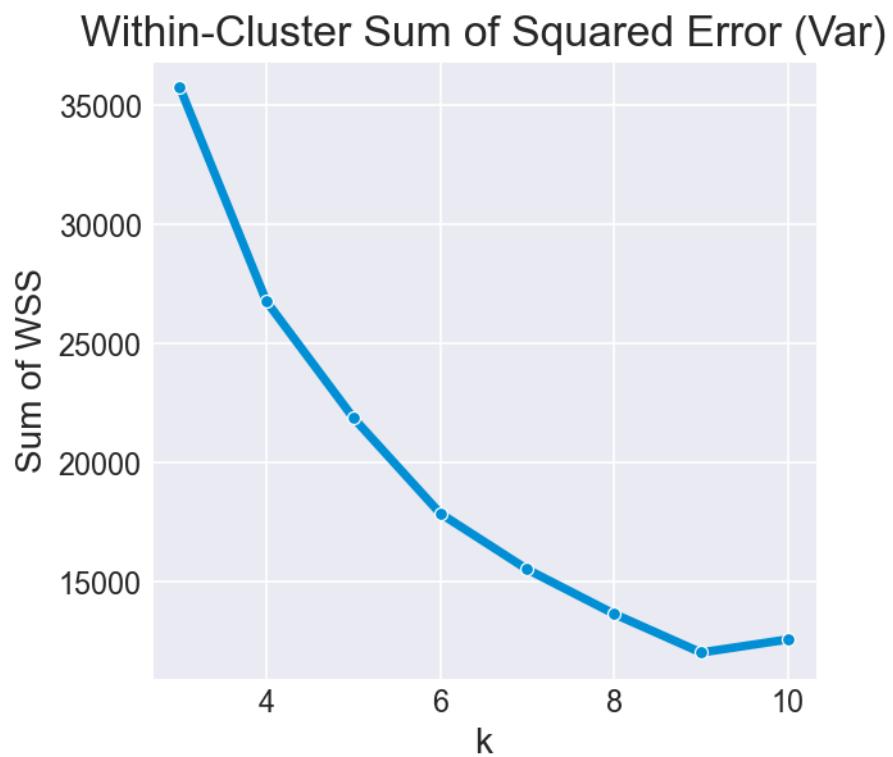


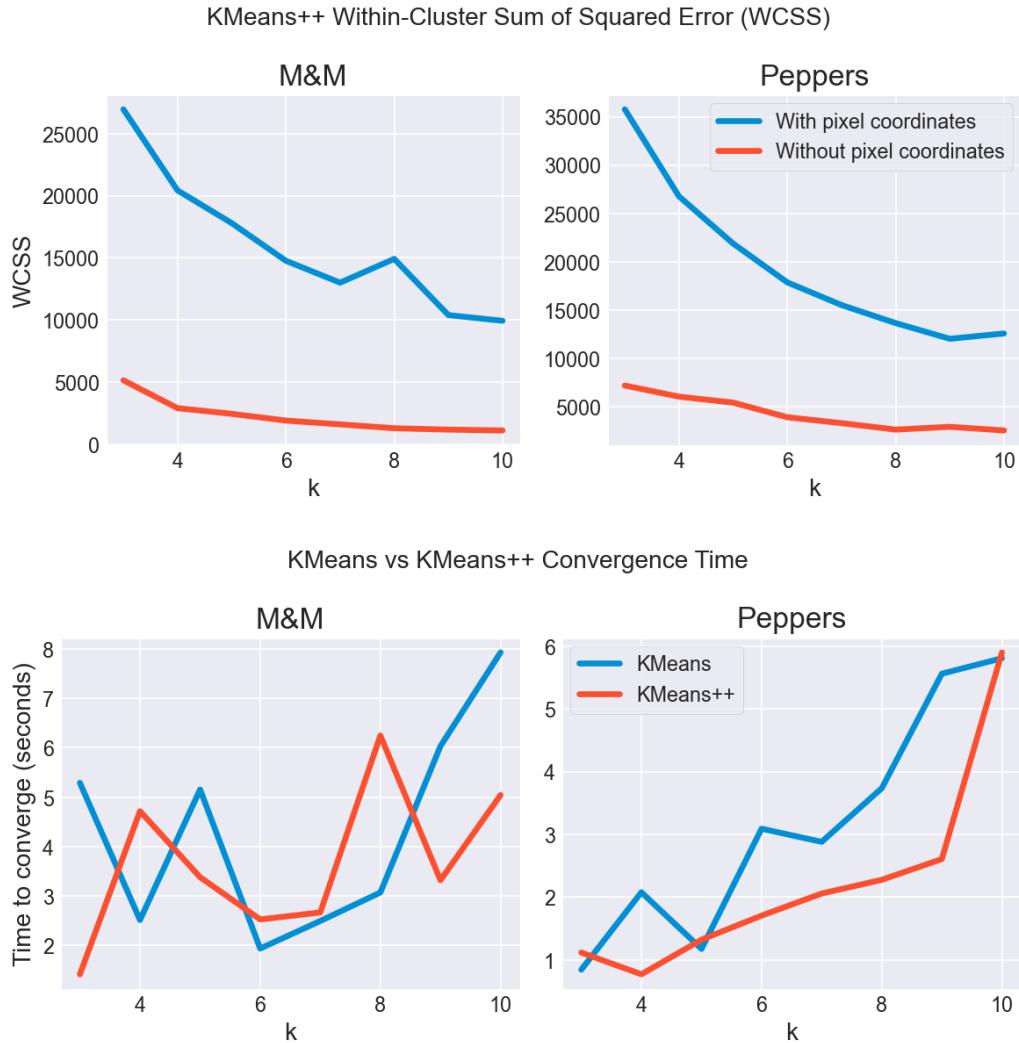
Image 2: Peppers



Using $k=8$



Model evaluation & comparison



Task 3: Eigenface

3.1 Explain why alignment is necessary for eigen-face

Alignment is necessary because face detection using eigenface is not invariant to position or scale. Each pixel coordinate position i, j is essentially treated as a distinct random variable (that may or may not correlate with others). If the images are not aligned, it will be difficult to find correlation or direction of maximum variance.

3.2. Perform PCA and show mean face

```
1 import numpy as np
2 from numpy import linalg as LA
3
4 class PCA:
5     def fit_transform(self, X: np.ndarray, n_components: int) → Tuple[np.ndarray, np.ndarray]:
6         """
7             input:
8                 X: n by m where n is the number of data points
9                 and m is the number of features
10            returns:
11                Xapprox: rank k approximation of X where k =
12                n_components
13
14            # create mean face by averaging the pixel
15            intensity values across all our flattened face images
16            self.mean = self.X.mean(0)
17            self.X_centered = X - self.mean
18
19            U, s, VT = LA.svd(X, full_matrices=False)
20            self.svd_output = (U, s, VT)
21            self.ev = self.explained_variance(s)
22
23            S = np.diag(s)
24            Xapprox = U[:, :n_components] @
25            S[:n_components, :n_components] @ VT[:n_components, :]
26            return Xapprox + self.mean
27
28        def explained_variance(self, s):
29            """calculate explained variance"""
30            ev = s ** 2 / np.sum(s ** 2)
31            return ev
32
33 k = 15
34 model = PCA()
35 A_approx = model.fit_transform(A, n_components=k)
```



Given the high dimensionality, directly performing eigen value decomposition of the covariance matrix would be slow. Find a faster way to compute eigen values and vectors, explain the reason.

Assumption: We don't care about eigenvalues that are approximately 0.

There are two ways (among other more complex methods) to compute eigenvalues and eigenvectors for non-square matrices with high dimensionality.

Let A be a real matrix $A \in \mathbb{R}^{n \times m}$ where $m > n$

1. Eigendecomposition

If A is a real matrix, then AA^T and $A^T A$ are symmetric square matrices.

Matrix $AA^T = L \in \mathbb{R}^{n \times n}$ and $A^T A = L \in \mathbb{R}^{m \times m}$.

Since the number of nonzero eigenvectors of AA^T is not greater than the number of eigenvectors of $A^T A$, we can perform eigendecomposition of $AA^T = L \in \mathbb{R}^{n \times n}$ to get eigenvectors and eigenvalues of $A^T A = L \in \mathbb{R}^{m \times m}$. Since $AA^T = L \in \mathbb{R}^{n \times n}$ is smaller (135×135) than ($45045, 45045$), it is much faster.

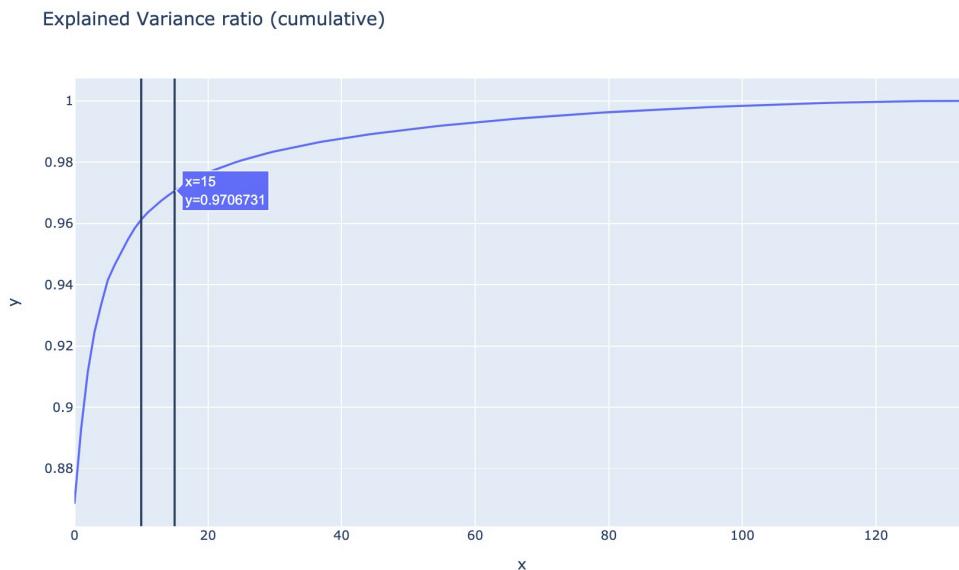
\therefore If v is an eigenvector with nonzero eigenvalue λ of $A^T A$, then Av is an eigenvector with the same eigenvalue of AA^T .

2. Singular Value Decomposition

Perform SVD directly on A , taking advantage of existing algorithms that can obtain eigenvalues and eigenvectors of the covariance matrix of A without having to compute $A^T A$.

3.3 Determine top k principal components and visualize top k eigenfaces

x axis represents the number of principal components, and y axis is the cumulative explained variance ratio.



First 15 principal components



3.4 Project test data to facespace spanned by first k eigenvectors and perform nearest neighbor search

```
idx = np.random.randint(0, test_images.shape[0])
x = test_images[idx]
x.shape
(45045,)

x_projected = x @ VT[:,k].T
x_projected.shape
(15,)

train_images_projected = A @ VT[:,k].T
train_images_projected.shape
(135, 15)

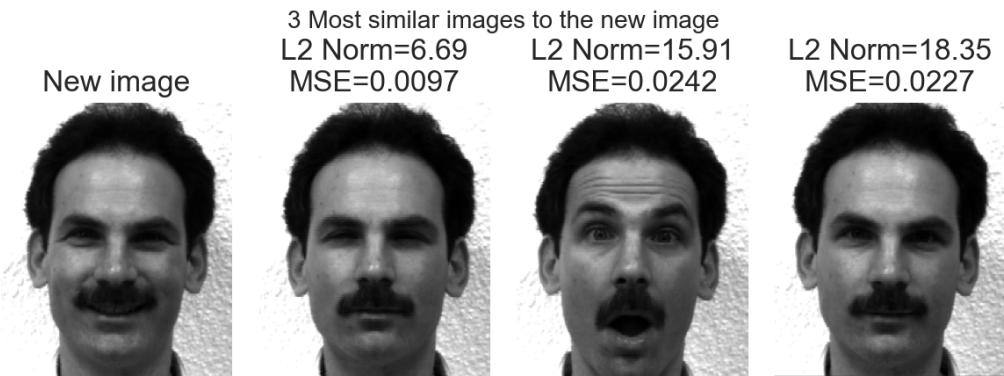
dist = np.sqrt(np.sum((train_images_projected - x_projected) ** 2, axis=1))
dist.shape
(135,)

plt.figure(figsize=(10, 4))
plt.suptitle('3 Most similar images to the new image')
nearest_idx = np.argsort(dist)

plt.subplot(1, 4, 1)
plt.imshow(x.reshape(dim), cmap='gray')
plt.grid(False)
plt.xticks([])
plt.yticks([])
plt.title('New image')

for i in range(3):
    idx = nearest_idx[i]
    mse = np.mean((x - A[idx]) ** 2)
    plt.subplot(1, 4, i+2)
    plt.imshow(A[idx].reshape(dim), cmap='gray')
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])
    plt.title(f'L2 Norm={dist[idx]:.2f}\nMSE={mse:.4f}')

plt.tight_layout()
print('MSE range [0, 1]')
```



3.5 Repeat the above step with a photo of self

```

selfies = np.array(selfies)
selfies = selfies.reshape(10, -1)
selfies.shape

(10, 45045)

idx = np.arange(selfies.shape[0], dtype=int)
np.random.shuffle(idx)
idx

array([7, 2, 5, 6, 4, 9, 8, 1, 0, 3])

selfies_test = selfies[idx[0]]
selfies_train = selfies[idx[1:]]

selfies_test_projected = selfies_test @ VT[:,k,:].T
print(selfies_test_projected.shape)

selfies_train_projected = selfies_train @ VT[:,k,:].T
print(selfies_train_projected.shape)

(15)
(9, 15)

dist = np.sqrt(np.sum((train_images_projected - selfies_test_projected) ** 2, axis=1))
nearest_idx = np.argsort(dist)

plt.figure(figsize=(10, 4))
plt.suptitle('3 Most similar images to new image')

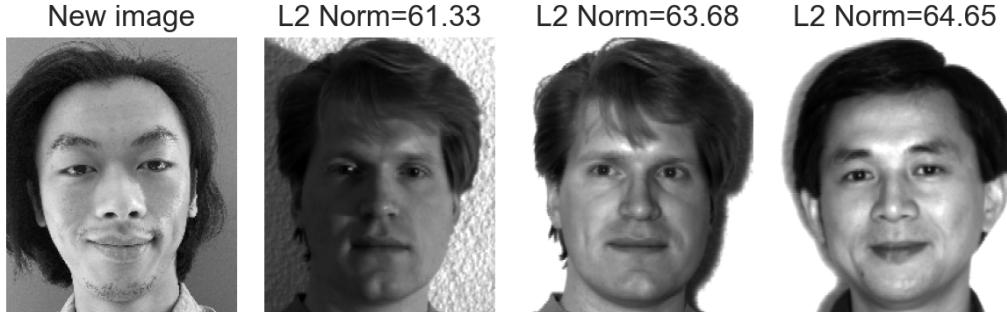
plt.subplot(1, 4, 1)
plt.imshow(selfies_test.reshape(dim), cmap='gray')
plt.grid(False)
plt.xticks([])
plt.yticks([])
plt.title('New image')

for i in range(3):
    idx = nearest_idx[i]
    plt.subplot(1, 4, i+2)
    plt.imshow(A[idx].reshape(dim), cmap='gray')
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])
    plt.title(f'L2 Norm={dist[idx]:.2f}')

plt.tight_layout()

```

3 Most similar images to new image



3.6 Repeat the above, with rest of selfies used as part of training set

```

train_images = A
combined_train_images_projected = np.vstack([train_images_projected, selfies_train_projected])
combined_train_images = np.vstack([train_images, selfies_train])

dist = np.sqrt(np.sum((combined_train_images_projected - selfies_test_projected) ** 2, axis=1))
nearest_idx = np.argsort(dist)

plt.figure(figsize=(10, 4))
plt.suptitle('3 Most similar images to new image')

plt.subplot(1, 4, 1)
plt.imshow(selfies_test.reshape(dim), cmap='gray')
plt.grid(False)
plt.xticks([])
plt.yticks([])
plt.title('New image')

for i in range(3):
    idx = nearest_idx[i]
    plt.subplot(1, 4, i+2)
    plt.imshow(combined_train_images[idx].reshape(dim), cmap='gray')
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])
    plt.title(f'L2 Norm={dist[idx]:.2f}')

plt.tight_layout()

```

3 Most similar images to new image

