

Kai Hirota

SID: u7233149

Task 1: 3D-2D Camera Calibration

1.1 Code for calibration

```
1 # helper functions not included: to_homogenous(), to_heterogenous()
2
3 def calibrate(im: np.ndarray, XYZ: np.ndarray, uv: np.ndarray):
4     """
5         Compute the 3x4 camera calibration matrix C such that xi = C @ Xi, where Xi is the
6         world coordinate and xi is the image pixel coordinate
7
8         Parameters:
9             im: Image of the calibration target.
10            XYZ: N x 3 array of XYZ coordinates of the calibration target points.
11            uv: N x 2 array of image coordinates of the calibration target points.
12
13         Returns:
14             C (np.ndarray): 3 x 4 camera calibration matrix
15             """
16         assert XYZ.shape[0] >= 6
17         assert uv.shape[0] == XYZ.shape[0]
18
19         n = uv.shape[0]
20
21         if uv.shape[1] == 2:
22             uv = to_homogenous_coord(uv)
23
24         if XYZ.shape[1] == 3:
25             XYZ = to_homogenous_coord(XYZ)
26
27         # construct matrix A and get the eigenvec corresponding to the smallest non-zero
28         # eigenvalue
29         A = get_A(uv, XYZ)
30         U, S, VT = LA.svd(A)
31         p = VT[-1]
```

```

30
31     # normalize p to unit length
32     p = p / (p @ p)
33     C = p.reshape(3, 4)
34
35     return C
36
37 def get_A(uv: np.ndarray, xyz: np.ndarray):
38     """Construct matrix A"""
39     n = uv.shape[0]
40     A = np.vstack([build_A(uv[i], xyz[i]) for i in range(n)])
41     assert A.shape == (2*n, 12)
42     return A
43
44 def build_A(img_coord, world_coord):
45     """Build A 2 rows at a time"""
46     if len(img_coord) == 3:
47         u, v, _ = img_coord
48     else:
49         u, v = img_coord
50
51     if len(world_coord) == 4:
52         x, y, z, _ = world_coord
53     else:
54         x, y, z = world_coord
55
56     # build matrix A of 2n x 12
57     A = np.array([
58         [x, y, z, 1, 0, 0, 0, -u*x, -u*y, -u*z, -u],
59         [0, 0, 0, 0, x, y, z, 1, -v*x, -v*y, -v*z, -v]
60     ])
61
62     return A
63
64 ##### Normalized DLT #####
65 h, w, _ = img.shape
66
67 Tnorm = np.array([
68     [w+h, 0, w/2],
69     [0, w+h, h/2],
70     [0, 0, 1]
71 ])
72
73 eigvals, eigvecs = LA.eig( (XYZ - XYZ.mean(0)) @ (XYZ - XYZ.mean(0)).T )
74 eigvals = np.diag(1 / eigvals[:3])
75 eigvecs = eigvecs[:3, :3]
76
77 left_component = eigvecs @ eigvals @ LA.inv(eigvecs)

```

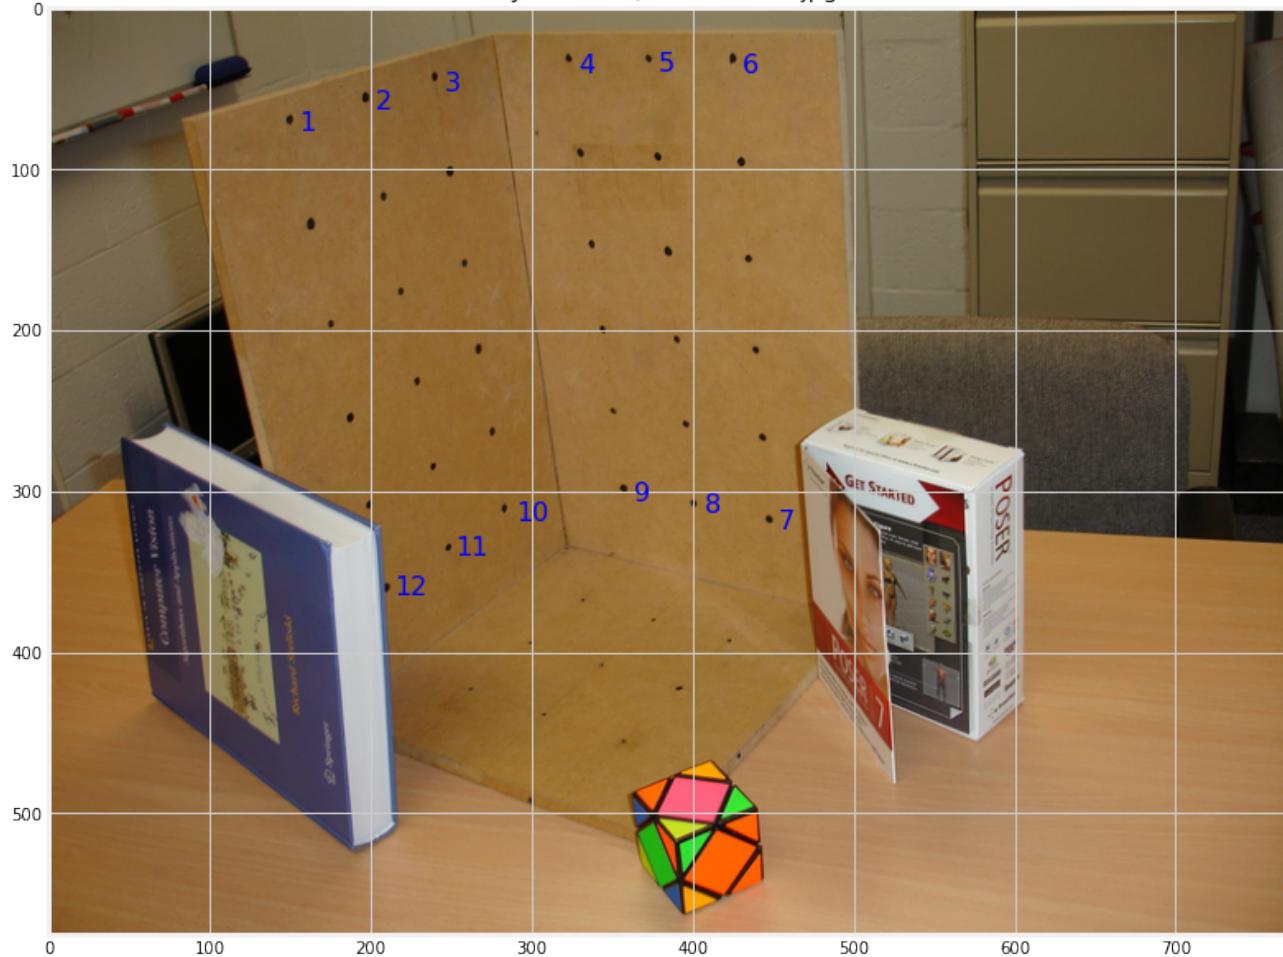
```

78 right_component = -eigvecs @ eigvals @ LA.inv(eigvecs) @ XYZ.mean(0)
79 right_component = right_component.reshape(len(right_component), -1)
80
81 Snorm = np.hstack([left_component, right_component])
82 Snorm = np.vstack([Snorm, [0,0,0,1]])
83
84 uv_norm = Tnorm @ to_homogenous_coord(uv).T
85 uv_norm = uv_norm.T
86
87 XYZ_norm = Snorm @ to_homogenous_coord(XYZ).T
88 XYZ_norm = XYZ_norm.T
89
90 #####
91
92 # how to use the code
93 P = calibrate(img, XYZ, uv)
94
95 # if using Normalized DLT
96 # P = calibrate(img, XYZ_norm, uv_norm)
97 # P = LA.inv(Tnorm) @ P @ Snorm
98 # P = P / P[-1,-1]
99
100 # project coordinates from world to img
101 xyz_proj = P @ XYZ.T
102 xyz_proj = to_heterogenous_coord(xyz_proj.T)

```

1.2 Selected Image & Correspondence Points

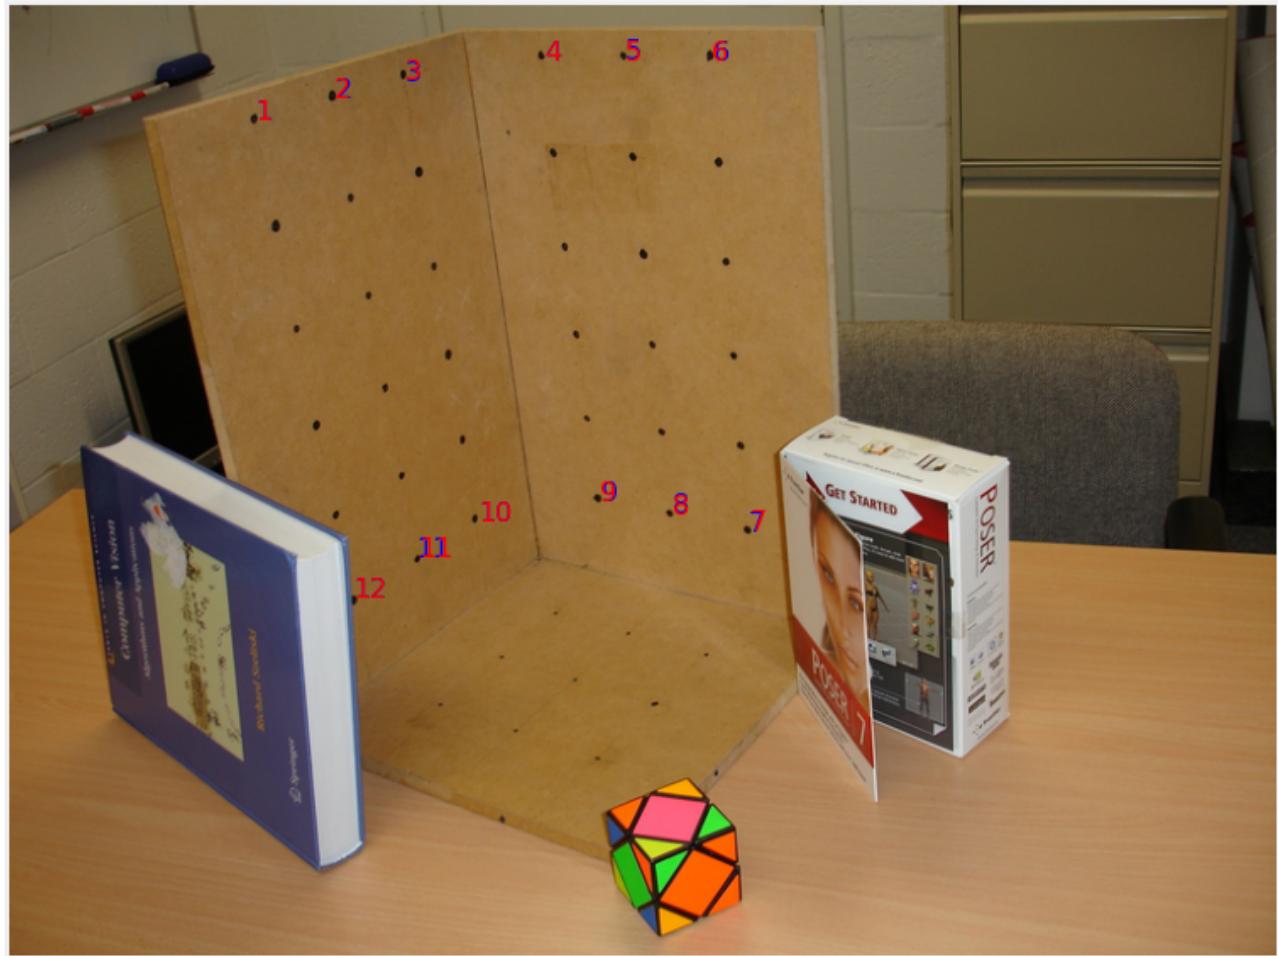
Python user/stereo2012a.jpg



1.3 Compute Calibration Matrix P and project correspondence points to pixel coordinate system

$$P = \begin{bmatrix} 0.0089 & -0.0041 & -0.0133 & 0.6928 \\ -0.0004 & -0.0156 & 0.0023 & 0.7208 \\ 0 & 0 & 0 & 0.0021 \end{bmatrix} \quad (1)$$

Original (blue) and projected (red) points



Mean Squared Error (pixels): 0.5426

1.4 Extrinsic and Intrinsic Parameters

$$K = \begin{bmatrix} 908.7838 & -1.4141 & 382.2355 \\ 0 & 892.3563 & 293.5467 \\ 0 & 0 & 1 \end{bmatrix} \quad (2)$$

$$R = \begin{bmatrix} 0.8174 & -0.1021 & -0.567 \\ 0.1568 & -0.9076 & 0.3894 \\ -0.5543 & -0.4072 & -0.7259 \end{bmatrix} \quad (3)$$

$$t = [76.4326 \quad 56.7049 \quad 85.5121] \quad (4)$$

1.5 Focal length & pitch angle

Focal Length

Since $K_{1,1} \neq K_{2,2}$ and $K_{1,2} \neq 0$, we have an intrinsic parameter matrix where the aspect ratio is not 1.

$$K = \begin{bmatrix} \alpha & \gamma & x_0 \\ 0 & \beta & y_0 \\ 0 & 0 & 1 \end{bmatrix} \quad (5)$$

where

- (x_0, y_0) = Principal point
- $\alpha = \alpha_x$
- $\beta = \frac{\alpha_y}{\sin \theta}$
- α_x, α_y are horizontal and vertical focal length in pixel units
- $\theta = \pi/2 \implies \beta = \alpha_y$

Assuming $\theta = \pi/2, \alpha_x = f_x$ and $\alpha_y = f_y$. Referring to K defined in (2) in [section 1.4](#), we have

- $f_x = 908.7838$ (pixels)
- $f_y = 892.3563$ (pixels)

Pitch Angle with respect to X-Z plane

The position, C , of the camera expressed in world coordinates is $C = -R^{-1}T = -R^T T$. Let \vec{p} be a vector that lie on the X-Z plane and \vec{c} be the position of the camera, with both vectors having the same basis (coordinate system). We have the geometric property of vectors:

$$\cos \theta = \frac{\vec{c} \cdot \vec{p}}{\|\vec{c}\| \cdot \|\vec{p}\|} \quad (6)$$

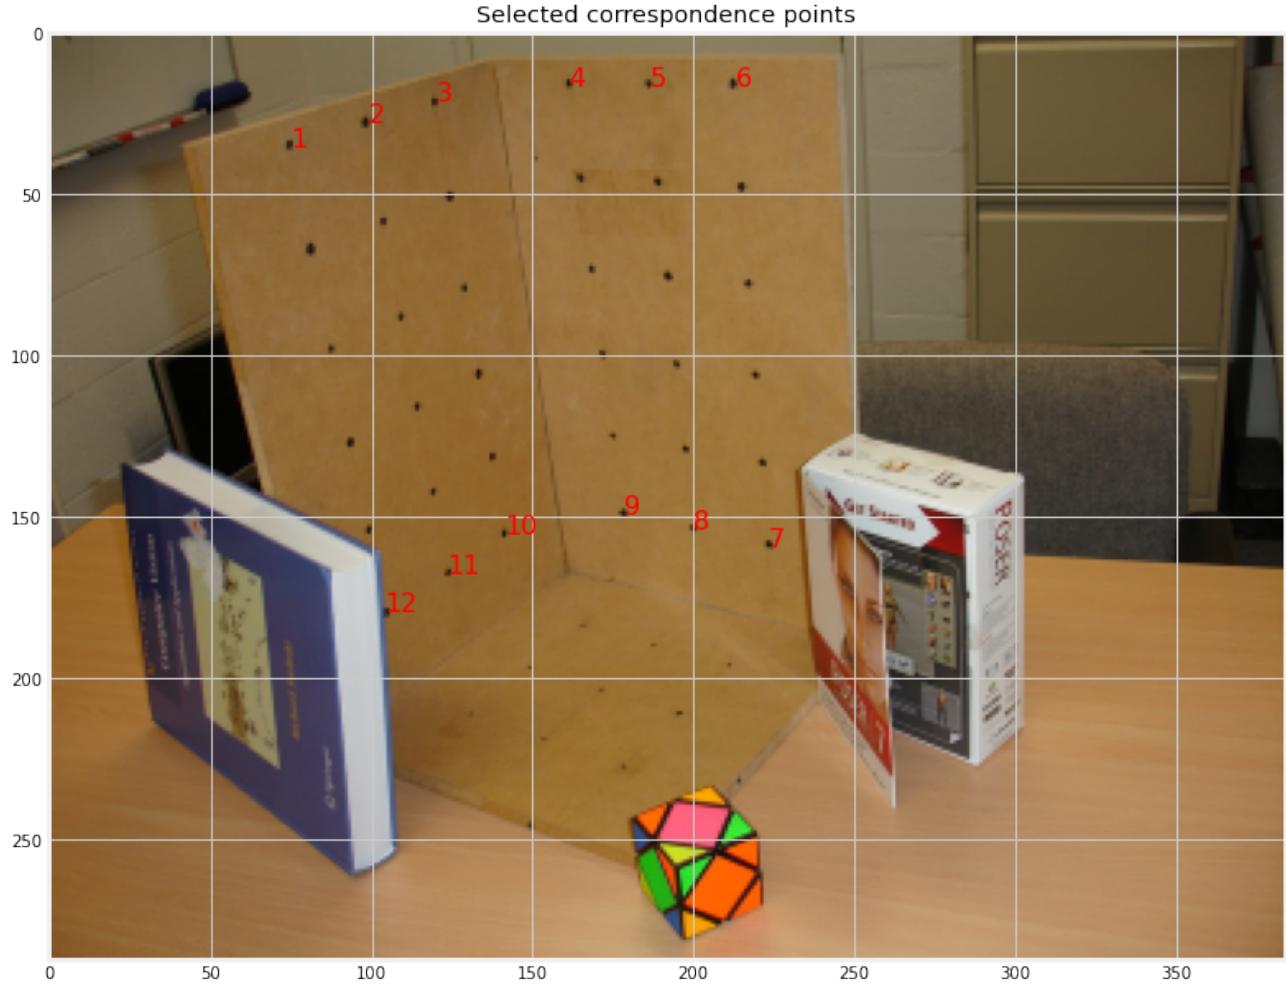
where $\vec{c}, \vec{p} \in \mathbb{R}^3$

The angle θ between C and the X-Z plane can be calculated as $\theta = \arccos(\cos(\theta))$.

```
1 | C = -R.T @ t.reshape(len(t), -1)
2 | C.T
3 | >> array([[-23.96402805,  94.08800595,  83.32400034]])
4 |
5 | plane = np.array([7,0,7])
6 | theta = np.arccos((C.T @ plane) / (LA.norm(C) * LA.norm(plane)))
7 | theta
8 | >> 1.23654328
9 |
10| np.degrees(theta)
11| >> 70.84871102
```

The pitch angle, or the angle between the camera's optical axis and the ground-plane, is approximately 70.85 degrees.

1.6 Resize



Resized image

Calibration Matrix P' and Parameters K', R', t'

$$P' = \begin{bmatrix} 0.0089 & -0.0041 & -0.0133 & 0.693 \\ -0.0002 & -0.0158 & 0.0026 & 0.7206 \\ 0.0 & 0.0 & 0.0 & 0.0043 \end{bmatrix} \quad (7)$$

$$K' = \begin{bmatrix} 442.1749 & 1.4406 & 182.1194 \\ 0.0 & 442.5788 & 135.2311 \\ 0.0 & 0.0 & 1.0 \end{bmatrix} \quad (8)$$

$$R' = \begin{bmatrix} 0.8133 & -0.1007 & -0.5731 \\ 0.1607 & -0.9077 & 0.3876 \\ -0.5592 & -0.4073 & -0.7221 \end{bmatrix} \quad (9)$$

$$t' = [73.4314 \quad 58.4466 \quad 83.2765] \quad (10)$$

- Intrinsic parameters K
 - Focal length $K'_{1,1}, K'_{2,2}$ are approximately half of $K_{1,1}, K_{2,2}$ due to width and height being halved.
 - Skew parameter $|K_{1,2}| \approx |K'_{1,2}|$ since the width and height were scaled by the same factor. In other

words, resizing while keeping the aspect ratio should not skew the image.

- Principal point $K'_{1,3}, K'_{2,3}$ are approximately half of $K_{1,3}, K_{2,3}$ due to the width and height being halved.
- Extrinsic parameters
 - $R \approx R'$ since rotation is scale-invariant if width and height are scaled by the same factor.
 - $t \approx t'$ for the same reason as R, R' above.

Task 2: Two-View DLT based homography estimation

2.1 Code, images used, and correspondence points

```
1 def homography(u2Trans: List[float], v2Trans: List[float],
2                 uBase: List[float], vBase: List[float]):
3     """
4         Computes the homography H applying the Direct Linear Transformation
5
6     Parameters:
7         u2Trans : Vectors with coordinates u and v of the transformed image
8         v2Trans : point (p')
9         uBase : vectors with coordinates u and v of the original base
10        vBase : image point p
11    """
12    uv_src = np.vstack([u2Trans, v2Trans]).T
13    uv_dest = np.vstack([uBase, vBase]).T
14    return _homography(uv_src, uv_dest)
15
16 def _homography(uv_src: np.ndarray, uv_dest: np.ndarray):
17     assert uv_src.shape[0] >= 6
18     assert uv_src.shape[0] == uv_dest.shape[0]
19
20     if uv_src.shape[1] == 2:
21         uv_src = to_homogenous_coord(uv_src)
22
23     if uv_dest.shape[1] == 2:
24         uv_dest = to_homogenous_coord(uv_dest)
25
26     A = get_A(uv_src, uv_dest)
27     U, S, VT = LA.svd(A)
28     p = VT[-1]
29
30     # normalize p to unit length
31     p = p / (p @ p)
32     H = p.reshape(3, 3)
33
34     return H
35
36 def get_A(uv_src: np.ndarray, uv_dest: np.ndarray):
```

```

37     """Iteratively build matrix A for homography"""
38     n = uv_src.shape[0]
39     A = np.vstack([build_A_homography(uv_src[i], uv_dest[i]) for i in range(n)])
40     assert A.shape == (2*n, 9)
41     return A
42
43 def build_A_homography(uv_src, uv_dest):
44     """Build A 2 rows at a time"""
45     assert len(uv_src) == 3 and len(uv_dest) == 3
46
47     u, v, _ = uv_src
48     x, y, _ = uv_dest
49
50     A = np.array([
51         [u, v, 1, 0, 0, 0, -x*u, -x*v, -x],
52         [0, 0, 0, u, v, 1, -y*u, -y*v, -y]
53     ])
54
55     return A
56
57 # how to use the code
58 # uv_left, uv_right: numpy array containing xy coordinates of correspondence points
59 H = homography(
60     u2Trans=uv_left[:, 0].tolist(),
61     v2Trans=uv_left[:, 1].tolist(),
62     uBase=uv_right[:, 0].tolist(),
63     vBase=uv_right[:, 1].tolist()
64 )
65
66 # warp the source image l using H
67 # l = Left.jpg = warp source
68 # r = Right.jpg = warp destination
69 x, y, _ = r.shape
70 size = (y, x)
71 l_warped = cv2.warpPerspective(l, H, dsize=size)
72
73 # project the correspondence points as well
74 uv_left = to_homogenous_coord(uv_left)
75 uv_left_proj = H @ uv_left.T
76 uv_left_proj = to_heterogenous_coord(uv_left_proj.T)

```

Selected correspondence points



2.2 Compute homography matrix H

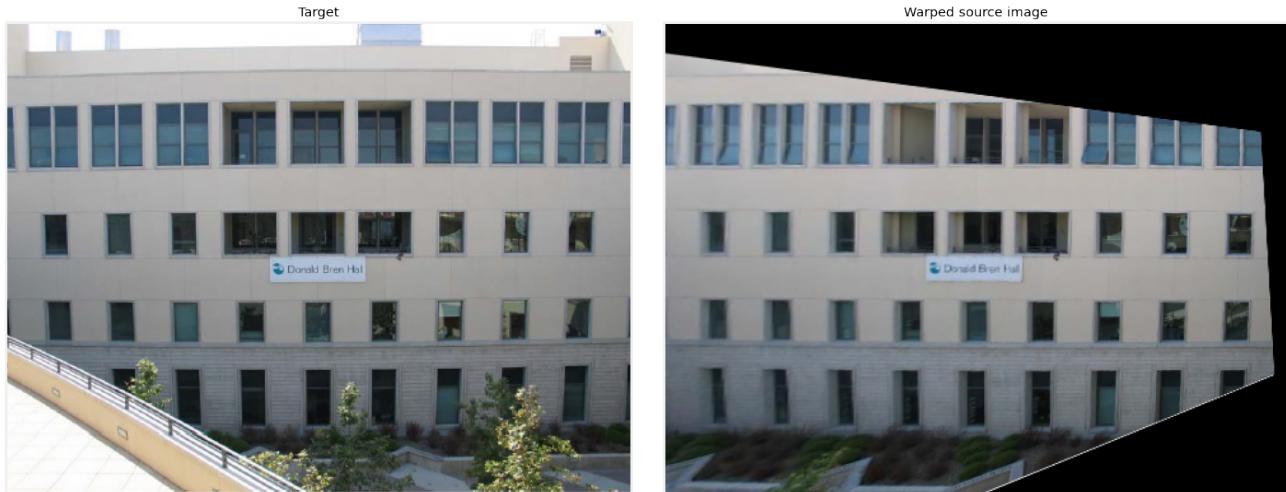
$$H = \begin{bmatrix} -0.0149 & 0.0004 & 0.9997 \\ -0.0024 & -0.0065 & 0.0169 \\ 0 & 0 & -0.0047 \end{bmatrix} \quad (11)$$

2.3 Warp image using H

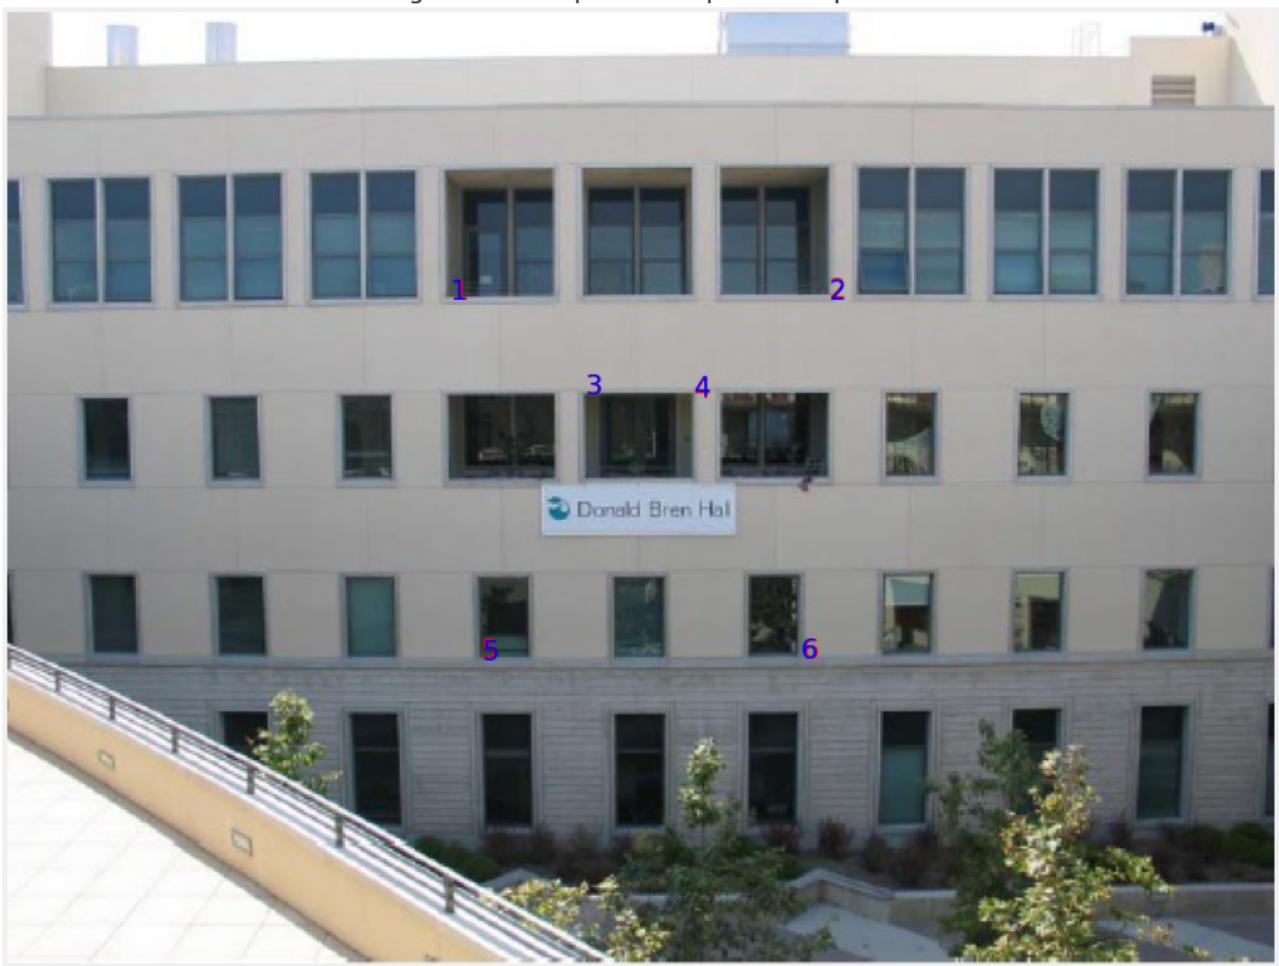
Result of warping using the computed Homography matrix



Comparison of warped (source) image with target image



Original and warped Correspondence points



```
1 | # euclidean distance between target and warped correspondence points in pixels
2 | dist = np.sqrt(np.sum(np.square(uv_right - uv_left_proj), 1))
3 | dist
4 | >> array([0.48353272, 0.59920041, 0.27162193, 1.05412312, 0.54916902,
5 |           0.2332485 ])
```

Mean Squared Error (pixels): 0.3556

The small distance between the target and warped points are result of manual input of correspondence points using `plt.ginput()`.